

### Question 1.

Tables and indexes:

Table name	Nb of records	Nb of pages in heap file (approximative is fine)	Index (type, field/column, alternative 1/2/3, clustered / unclustered)
ARTIST	1000000	6396	NORMAL, LASTNAME, alternative 2, unclustered
			NORMAL, IDARTIST, alternative 2, unclustered
MOVIE1	449833	13821	NORMAL, IDMOVIE, alternative 2, unclustered
MOVIE2	449833	13821	NORMAL, IDMOVIE, alternative 2, unclustered
			NORMAL, YEAR, alternative 2, unclustered
			NORMAL, CODECOUNTRY, alternative 2, unclustered
MOVIE3	449833	13821	NORMAL, IDMOVIE, alternative 2, unclustered
			BITMAP, YEAR, alternative 2, unclustered
			BITMAP, CODECOUNTRY, alternative 2, unclustered
MOVIE4	449833	null	IOT-TOP, IDMOVIE, alternative 1, clustered
			NORMAL, CODECOUNTRY, alternative 1, clustered
MOVIERATER	500000	6045	NORMAL, EMAIL, alternative 2, unclustered
RATING	629461	5224	NORMAL, IDMOVIE EMAIL, alternative 2, unclustered
COUNTRY	180	2	NORMAL, CODE, alternative 2, unclustered
ROLE	539981	2285	NORMAL, IDMOVIE IDACTOR, alternative 2, unclustered

### Question 2. Fill the following tables:

Attribute name	Cardinality (number of distinct values)	Minimal value	Maximal value
MOVIE1.IDMOVIE	449833	0	500000
MOVIE1.YEAR	20	1991	2010
MOVIE1.CODECOUNTRY	180	aaau	aahr

### Question 3.1 On MOVIE1 and MOVIE4 :

SELECT TITLE FROM MOVIE<sub>i</sub> WHERE IDMOVIE=50273 ;

TABLE	MODE	OPERATORS	BUFFER READS	DISK READS
MOVIE1	CHOOSE	INDEX UNIQUE SCAN + BY INDEX ROWID	4	4
MOVIE4	CHOOSE	INDEX UNIQUE SCAN	3	3

Your explanation of the SQL query (in plain English) and of the execution plan:  
This query gives the title of movie 50273.

In the form above, in CHOOSE mode we use the index, since it is there. Since MOVIE1 use index of Alternative 2, it need to read the index pages and read the responding heap pages. But MOVIE4 use index of Alternative 1, it only need to read the index pages.

Question 3.2 On MOVIE2 and MOVIE4 :

SELECT COUNT(\*) FROM MOVIEi WHERE IDMOVIE BETWEEN 55273 AND 60000 ;

TABLE	MODE	OPERATORS	BUFFER READS	DISK READS
<i>MOVIE2</i>	<i>CHOOSE</i>	INDEX RANGE SCAN	12	12
<i>MOVIE4</i>	<i>CHOOSE</i>	INDEX RANGE SCAN	168	168

Your explanation of the SQL query (in plain English) and of the execution plan:  
This query gives the number of movies between 55273 and 60000, which is 4215.  
In the form above, in CHOOSE mode we use the index, since range conditions are sensitive to clustering. MOVIE4 has more leaf blocks than MOVIE2, so index range scan in MOVIE4 needs to read more pages than in MOVIE2.

Question 3.3 On MOVIE2, with RULE, CHOOSE :

SELECT TITLE FROM MOVIEi WHERE YEAR=1999 ;

TABLE	MODE	OPERATORS	BUFFER READS	DISK READS
<i>MOVIE2</i>	<i>RULE</i>	INDEX RANGE SCAN + BY INDEX ROWID	11899	11247
<i>MOVIE2</i>	<i>CHOOSE</i>	TABLE ACCESS FULL	14259	13822

Your explanation of the SQL query (in plain English) and of the execution plan:  
This query gives the title of movies, which year is 1999.  
In the form above, in RULE mode we use the index, since it is there.  
However, in CHOOSE mode the optimizer understands that using the index would be too expensive, thus a full scan is done instead.

Question 3.4 On MOVIE2, MOVIE3, MOVIE4 :

SELECT COUNT(\*) FROM MOVIEi WHERE YEAR=1999 ;

TABLE	MODE	OPERATORS	BUFFER READS	DISK READS
<i>MOVIE2</i>	<i>CHOOSE</i>	INDEX RANGE SCAN	55	55
<i>MOVIE3</i>	<i>CHOOSE</i>	BITMAP INDEX SINGLE VALUE	8	8
<i>MOVIE4</i>	<i>CHOOSE</i>	INDEX FAST FULL SCAN	17743	17723

Your explanation of the SQL query (in plain English) and of the execution plan:  
This query gives the number of movies, which year is 1999.  
In the form above, in CHOOSE mode we use the index, since MOVIE2, MOVIE3 have index on YEAR, since we only need to count the number, finding the corresponding page in heap file is unnecessary. Each bit in the bitmap corresponds to a rowId, so MOVIE3 requires the least pages to read, MOVIE2 is the second. MOVIE4 has no index on YEAR, so it full scans the table, which needs to read a lot of pages.

Question 3.5 On MOVIE2, MOVIE3, with RULE, CHOOSE :

SELECT TITLE FROM MOVIEi WHERE CODECOUNTRY='aaej' ;

TABLE	MODE	OPERATORS	BUFFER READS	DISK READS
MOVIE2	RULE	INDEX RANGE SCAN + BY INDEX ROWID	2377	2323
MOVIE2	CHOOSE	INDEX RANGE SCAN + BY INDEX ROWID	2377	2271
MOVIE3	RULE	TABLE ACCESS FULL	13875	13822
MOVIE3	CHOOSE	BITMAP INDEX SINGLE VALUE + BY INDEX ROWID	2322	2317

Your explanation of the SQL query (in plain English) and of the execution plan:  
 This query gives the title of movies, which CODECOUNTRY is 'aaej'.  
 In the MOVIE2, in RULE and CHOOSE mode we use the index, since it has index on CODECOUNTRY. In the MOVIE3, in RULE mode we use full scan since it is easy and cheap. In CHOOSE mode, the optimizer find it has bitmap index on CODECOUNTRY, thus a bitmap index is done instead.

## SECTION 4 – MULTI-TABLE (JOIN) QUERIES

Question 4.1 On MOVIE2 and MOVIE3, with RULE, CHOOSE :

SELECT TITLE, LASTNAME FROM MOVIEi M, ARTIST A WHERE  
 CODECOUNTRY='aaej' AND M.IDMES=A.IDARTIST ;

TABLE	MODE	OPERATORS	BUFFER READS	DISK READS
MOVIE2	CHOOSE	INDEX RANGE SCAN + BY INDEX ROWID	2323	2323
MOVIE2	RULE	INDEX RANGE SCAN + BY INDEX ROWID	2377	2323
MOVIE3	CHOOSE	BITMAP INDEX SINGLE VALUE + BY INDEX ROWID	2317	2317
MOVIE3	RULE	TABLE ACCESS FULL	13875	13822

Your explanation of the SQL query (in plain English) and of the execution plan:  
 This query gives the title of the movie, which have been directed by an actor and CODECOUNTRY equals 'aaej', and the lastname of the corresponding director.  
 In the MOVIE2, in RULE and CHOOSE mode we use the index, since it has index on CODECOUNTRY. After selecting the codecounrty, MOVIE2 uses HASH JOIN to join ARTIST.  
 In the MOVIE3, in RULE mode we use full scan since it is easy and cheap. In CHOOSE mode, the optimizer find it has bitmap index on CODECOUNTRY, thus a bitmap index is done instead. After selecting the codecounrty, MOVIE3 uses HASH JOIN to join ARTIST.

Question 4.2 On MOVIE2 :

SELECT lastname, title  
 FROM MOVIEi M, MOVIERATER MR, RATING R  
 WHERE R.email='aaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaa' AND  
 MR.email= R.email AND M.IDMOVIE=R.IDMOVIE;

TABLE	MODE	OPERATORS	BUFFER READS	DISK READS
-------	------	-----------	--------------	------------

<i>MOVIE2</i>	<i>CHOOSE</i>	INDEX UNIQUE SCAN + BY INDEX ROWID	3	3
---------------	---------------	---------------------------------------	---	---

Your explanation of the SQL query (in plain English) and of the execution plan:  
This query gives the title of movies and lastname of the movie rater, where the movie's movie rater's email is 'aaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaa'. The result is 'nfs knvi' 'owzd'.  
In CHOOSE mode we use the index. After joining the MOVIERATER and RATING through NESTED LOOPS, we can get the IDMOVIE. Since MOVIE2 has index on IDMOVIE, we use index unique scan to read the pages when join through NESTED LOOPS.

Question 4.3 On MOVIE1 :

```
SELECT lastname, title
FROM MOVIEi M, MOVIERATER MR, RATING R
WHERE MR.email= R.email AND M.IDMOVIE=R.IDMOVIE AND R.IDMOVIE=
367856;
```

TABLE	MODE	OPERATORS	BUFFER READS	DISK READS
<i>MOVIE1</i>	<i>CHOOSE</i>	INDEX UNIQUE SCAN + BY INDEX ROWID	4	4

Your explanation of the SQL query (in plain English) and of the execution plan:  
This query gives the title of movies and lastname of movie rater, where the movie id is 367856. The result is 'nfs knvi' 'owzd' and 'pp' 'owzd'.  
In CHOOSE mode we use the index. Since MOVIE1 has index on IDMOVIE, we first select IDMOVIE=367856 by index, then using NESTED LOOPS JOIN to join other tables.

Question 4.4 On MOVIE3 :

```
SELECT LASTNAME, COUNT(*)
FROM MOVIEi M, ARTIST A
WHERE YEAR=1999 AND M.IDMES=A.IDARTIST
GROUP BY A.LASTNAME;
```

TABLE	MODE	OPERATORS	BUFFER READS	DISK READS
<i>MOVIE3</i>	<i>CHOOSE</i>	BITMAP INDEX SINGLE VALUE + BY INDEX ROWID	11200	11200

Your explanation of the SQL query (in plain English) and of the execution plan:  
This query gives lastname of the director of movies that were released in 1999 and were directed by an actor, and the number of films directed by actors and released in 1999 according to lastname.  
In CHOOSE mode we use the bitmap index. Since MOVIE3 has bitmap index on YEAR, we first select YEAR=1999 by index, then using HASH JOIN to join other tables.

Question 4.5 On MOVIE3 :

```
SELECT LASTNAME, COUNT(*)
FROM ARTIST
WHERE IDARTIST IN (SELECT DISTINCT IDMES
                    FROM MOVIEi
                    WHERE YEAR=1999)
GROUP BY LASTNAME;
```

TABLE	MODE	OPERATORS	BUFFER READS	DISK READS
MOVIE3	CHOOSE	BITMAP INDEX SINGLE VALUE + BY INDEX ROWID	11200	11200

Your explanation of the SQL query (in plain English) and of the execution plan:  
This query gives lastname of the director of movies that were released in 1999 and were directed by an actor, and the number of actors who directed the movie and released in 1999 according to lastname.  
In CHOOSE mode we use the bitmap index. Since MOVIE3 has bitmap index on YEAR, we first select YEAR=1999 by index, then using HASH JOIN to join other tables.

Question 4.6 On Movie3 :

```
SELECT COUNT(*)
FROM MOVIEi M, ROLE R, ARTIST A
WHERE M.IDMOVIE=R.IDMOVIE AND R.IDACTOR=A.IDARTIST AND
      M.IDMES=A.IDARTIST AND CODECOUNTRY='aej' AND YEAR=1999;
```

TABLE	MODE	OPERATORS	BUFFER READS	DISK READS
MOVIE3	CHOOSE	BITMAP INDEX SINGLE VALUE + BY INDEX ROWID	167	166

Your explanation of the SQL query (in plain English) and of the execution plan:  
This query gives the number of movies which was released in 1999 with CODECOUNTRY 'aej' and have been directed by one of their actors.  
In CHOOSE mode we use the bitmap index. Since MOVIE3 has bitmap index on YEAR and CODECOUNTRY, we first select YEAR=1999 and CODECOUNTRY='aej' by index, then using NESTED LOOPS JOIN to join other tables.

Question 4.7 On MOVIE2 and MOVIE4 :

```
SELECT TITLE
FROM MOVIEi M, RATING R
WHERE M.IDMOVIE = R.IDMOVIE
GROUP BY TITLE
HAVING AVG(RATE) > 15;
```

TABLE	MODE	OPERATORS	BUFFER READS	DISK READS
MOVIE2	CHOOSE	TABLE ACCESS FULL	13824	13822

<i>MOVIE4</i>	<i>CHOOSE</i>	INDEX FAST FULL SCAN	17743	17723
---------------	---------------	----------------------	-------	-------

Your explanation of the SQL query (in plain English) and of the execution plan:  
This query gives title of movies, which average rate of same titles is higher than 15.  
In CHOOSE mode, we use the FULL SCAN in MOVIE2 and MOVIE4, the optimizer understands that there is no filtering before and using the index would be too expensive, thus a full scan is done instead.

Question 4.8 On MOVIE2 and MOVIE4 :

```

SELECT TITLE
FROM MOVIEi M, (SELECT IDMOVIE, AVG(RATE) as AVG_RATE
FROM RATING
GROUP BY IDMOVIE) N
WHERE M.IDMOVIE = N.IDMOVIE AND AVG_RATE > 15;

```

TABLE	MODE	OPERATORS	BUFFER READS	DISK READS
<i>MOVIE2</i>	<i>CHOOSE</i>	TABLE ACCESS FULL	14977	13822
<i>MOVIE4</i>	<i>CHOOSE</i>	INDEX FAST FULL SCAN	18874	17723

Your explanation of the SQL query (in plain English) and of the execution plan:  
This query gives title of movies, which average rate of same IDMOVIE is higher than 15.  
In CHOOSE mode, we use the FULL SCAN in MOVIE2 and MOVIE4, the optimizer understands that using the index would be too expensive, thus a full scan is done instead.